super_block 구조체 include/linux/fs.h dentry 구조체 include/linux/dcache.h → struct super_block { struct dentry { struct dentry *s_root; -→ struct dentry { struct hlist_bl_heads_anon; → struct hlist_bl_node d_hash; → struct hlist_bl_node d_hash; struct super_block *d_sb; struct inode *d_inode; _ struct dentry *d_parent; task_struct 구조체 union { ----- struct inode *d_inode; path 구조체 include/linux/sched.h struct lockref d_lockref; → struct hlist_node d_alias;

 include/linux/path.h files_struct 구조체 fs/dcache.c struct workqueue_struct *s_dio_done_wq; } d_u; include/linux/fdtable.h struct task_struct { static struct hlist_bl_head *dentry_hashtable; struct vfsmount *mnt; struct list_head d_lru; struct files_struct *files; struct user_namespace *s_user_ns; struct dentry *dentry; _____ struct files_struct { wait_queue_head_t *d_wait; struct fdtable fdtab; struct list_head d_child; file 구조체 struct rcu_head rcu; struct list_head d_subdirs; include/linux/fs.h struct work_struct destroy_work; struct file { struct list_head s_inodes; ◀ ____ struct hlist_node d_alias; _ struct path struct list_head s_inodes_wb; ← ___ struct inode *f_inode; struct hlist_bl_node d_in_lookup_hash; __ struct address_space *f_mapping; struct rcu_head d_rcu; *private_data; struct hlist_bl_node d_hash; → struct list_head d_lru; ◆ struct dentry { →struct hlist_bl_node d_hash; → struct list_head d_lru; } d_u; inode 구조체 (Regular file) inode 구조체 (bdev file) include/linux/fs.h include/linux/fs.h struct inode { struct inode { struct super_block *i_sb; struct super_block *i_sb; struct address_space *i_mapping; ___ struct address_space *i_mapping; ———— → struct hlist_node i_hash; → struct hlist_node i_hash;
→ struct list_head i_io_list; struct list_head i_io_list; struct list_head i_lru; struct list_head i_lru; struct list_head i_sb_list; → struct list_head i_sb_list; include/linux/mmzone.h struct list_head i_wb_list; struct list_head i_wb_list; struct per_cpu_pageset { — struct hlist_head i_dentry; ◄ struct per_cpu_pages { struct rcu_head i_rcu; struct rcu_head i_rcu; struct list_head lists[MIGRATE_PCPTYPES]; } pcp; fs/inode.c struct address_space { struct address_space { ◀ struct inode *host; ____ struct inode *host; static struct hlist_head *inode_hashtable; *private_data; void *private_data; }i_data; }i_data; arch/x86/include/asm/mmzone_64.h include/linux/mmzone.h *i_private; *i_private; extern struct pglist_data *node_data[]; void void typedef struct pglist_data { struct pglist_data *zone_pgdat; struct per_cpu_pageset __percpu *pageset; struct inode { ----- struct free_area { struct inode { → struct hlist_node i_hash;
→ …… → struct hlist_node i_hash;
→ struct list_head free_list[MIGRATE_TYPES]; → struct hlist_node i_hash; } free_area[MAX_ORDER]; • • • → struct list_head i_lru;
→ struct list_head i_lru; } node_zones[MAX_NR_ZONES]; struct address_space { struct inode *host; struct zonelist { struct radix_tree_root { struct zoneref { struct inode { struct radix_tree_node __rcu *rnode; _____ struct radix_tree_node { _____ struct zone *zone; → struct hlist_node i_hash; } _zonerefs[MAX_ZONES_PER_ZONELIST + 1]; struct radix_tree_node *parent; — → NULL } node_zonelists[MAX_ZONELISTS]; → struct list_head i_lru; struct radix_tree_root *root; struct rb_root i_mmap; void __rcu *slots[RADIX_TREE_MAP_SIZE]; /* means !SPARSEMEM */ nrpages; #ifdef CONFIG_FLAT_NODE_MEM_MAP writeback_index; Physical address Physical address ____ struct page *node_mem_map; _____struct list_lru s_inode_lru ____cacheline_aligned_in_smp;ፈ struct list_head private_list; *private_data; Node #0 struct mount { struct list_head s_mounts; struct list_head mnt_instance; free_area[MAX_ORDER]

free page struct mount *mnt_parent; // Order of buddy allocator struct vfsmount { Page struct array struct zone node_zones[] Single node struct dentry *mnt_root; — — struct super_block *mnt_sb; → struct list_head mnt_mounts; — [ZONE_DMA] -----Node #n struct list_head mnt_child; struct mnt_namespace { struct task_struct { — struct mtd_info *s_mtd; struct dentry *mnt_mountpoint; struct nsproxy *nsproxy; U struct mnt_namespace *mnt_ns; U struct list_head list; struct list_head mnt_list; -'-------— struct mnt_namespace *mnt_ns; — struct backing_dev_info *s_bdi; DISCONTIGMEM **FLATMEM** — struct block_device*s_bdev; [ZONE_DMA32] *s_fs_info; -→ struct xxx_sb_info { // File system private info. ┌**〈►** struct mount { ___ struct mount *mnt_parent; fs/super.c struct dentry *mnt_mountpoint; static struct list_head super_blocks; struct vfsmount { struct super_block { struct super_block { struct dentry *mnt_root; — → struct list_head s_list; struct list_head s_list; —— struct list_head s_list; struct super_block *mnt_sb; · <u>Ł</u> per_cpu_pageset
// caching 1 pages → struct hlist_node s_instances;
→ … struct hlist_node s_instances; struct hlist_node s_instances; ◄— struct list_head mnt_child; -----[ZONE_NORMAL] struct list_head mnt_mounts; struct list_head mnt_instance; fs/filesystems.c struct list_head mnt_list; static struct file_system_type *file_systems ___ struct mnt_namespace *mnt_ns; ______ '---<u>-</u>--file_system_type 구조체 include/linux/fs.h struct file_system_type { struct file_system_type { ┌──► NULL const char *name; const char *name; int fs_flags; int fs_flags; struct dentry *(*mount) (struct struct dentry *(*mount) (struct struct mount { ___ struct mount *mnt_parent; void (*kill_sb) (struct super_block *); void (*kill_sb) (struct super_block *); struct dentry *mnt_mountpoint; struct module *owner; struct module *owner; struct file_system_type * next; ---struct file_system_type * next; ——— struct vfsmount { → struct hlist_head fs_supers; struct dentry *mnt_root; struct super_block *mnt_sb; — struct list_head mnt_child; <---</p> struct list_head mnt_mounts; struct list_head mnt_instance; struct block_device { struct mtd_info { — struct list_head mnt_list; ◄ ___ struct mnt_namespace *mnt_ns; * A Memory Technology Device
* is a type of device file in Linux
* for interacting with flash memory. — struct super_block *bd_super; bdev_inode 구조체 fs/block_dev.c — struct block_device * bd_contains; fs/block_dev.c — struct backing_dev_info *bd_bdi; static struct list_head all_bdevs; struct bdev_inode { — struct request_queue * bd_queue; struct block_device { → struct list_head bd_list;
→ struct list_head bd_list;
→ struct list_head bd_list; _____ struct gendisk * bd_disk; struct hlist_node i_hash; ◀ struct inode *host; struct gendisk { struct address_space *i_mapping; struct request_queue *queue;

> union { ------ struct block_device *i_bdev;

}vfs_inode;

struct request_queue {

struct backing_dev_info *backing_dev_info;